

## Import necessary libraries

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor
from xgboost import XGBRegressor
```

## Load the dataset

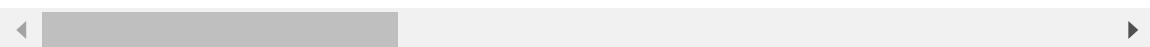
```
In [4]: data = pd.read_csv('C:/Users/91700/OneDrive/Desktop/SARTHAK/Project UM/Data/laptops.csv')
```

## Information Related to Data

```
In [6]: data.head()
```

	Company	Product	TypeName	Inches	Ram	OS	Weight	Price_euros	Screen
0	Apple	MacBook Pro	Ultrabook	13.3	8	macOS	1.37	1339.69	Standard
1	Apple	Macbook Air	Ultrabook	13.3	8	macOS	1.34	898.94	Standard
2	HP	250 G6	Notebook	15.6	8	No OS	1.86	575.00	Full HD
3	Apple	MacBook Pro	Ultrabook	15.4	16	macOS	1.83	2537.45	Standard
4	Apple	MacBook Pro	Ultrabook	13.3	8	macOS	1.37	1803.60	Standard

5 rows × 23 columns



```
In [7]: data.shape
```

```
Out[7]: (1275, 23)
```

In [8]: `data.describe()`

Out[8]:

	Inches	Ram	Weight	Price_euros	ScreenW	ScreenH	
<b>count</b>	1275.000000	1275.000000	1275.000000	1275.000000	1275.000000	1275.000000	1
<b>mean</b>	15.022902	8.440784	2.040525	1134.969059	1900.043922	1073.904314	
<b>std</b>	1.429470	5.097809	0.669196	700.752504	493.346186	283.883940	
<b>min</b>	10.100000	2.000000	0.690000	174.000000	1366.000000	768.000000	
<b>25%</b>	14.000000	4.000000	1.500000	609.000000	1920.000000	1080.000000	
<b>50%</b>	15.600000	8.000000	2.040000	989.000000	1920.000000	1080.000000	
<b>75%</b>	15.600000	8.000000	2.310000	1496.500000	1920.000000	1080.000000	
<b>max</b>	18.400000	64.000000	4.700000	6099.000000	3840.000000	2160.000000	

◀ ▶

In [9]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1275 entries, 0 to 1274
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Company          1275 non-null   object 
 1   Product          1275 non-null   object 
 2   TypeName         1275 non-null   object 
 3   Inches           1275 non-null   float64
 4   Ram              1275 non-null   int64  
 5   OS               1275 non-null   object 
 6   Weight           1275 non-null   float64
 7   Price_euros     1275 non-null   float64
 8   Screen            1275 non-null   object 
 9   ScreenW          1275 non-null   int64  
 10  ScreenH          1275 non-null   int64  
 11  Touchscreen      1275 non-null   object 
 12  IPSpanel         1275 non-null   object 
 13  RetinaDisplay    1275 non-null   object 
 14  CPU_company      1275 non-null   object 
 15  CPU_freq          1275 non-null   float64
 16  CPU_model         1275 non-null   object 
 17  PrimaryStorage    1275 non-null   int64  
 18  SecondaryStorage  1275 non-null   int64  
 19  PrimaryStorageType 1275 non-null   object 
 20  SecondaryStorageType 1275 non-null   object 
 21  GPU_company       1275 non-null   object 
 22  GPU_model          1275 non-null   object 

dtypes: float64(4), int64(5), object(14)
memory usage: 229.2+ KB
```

In [10]: `data.isnull().sum()`

```
Out[10]: Company          0
          Product         0
          TypeName        0
          Inches          0
          Ram             0
          OS              0
          Weight          0
          Price_euros     0
          Screen          0
          ScreenW         0
          ScreenH         0
          Touchscreen     0
          IPSpanel        0
          RetinaDisplay   0
          CPU_company     0
          CPU_freq         0
          CPU_model        0
          PrimaryStorage   0
          SecondaryStorage 0
          PrimaryStorageType 0
          SecondaryStorageType 0
          GPU_company      0
          GPU_model        0
          dtype: int64
```

```
In [11]: data.corr(numeric_only = True)
```

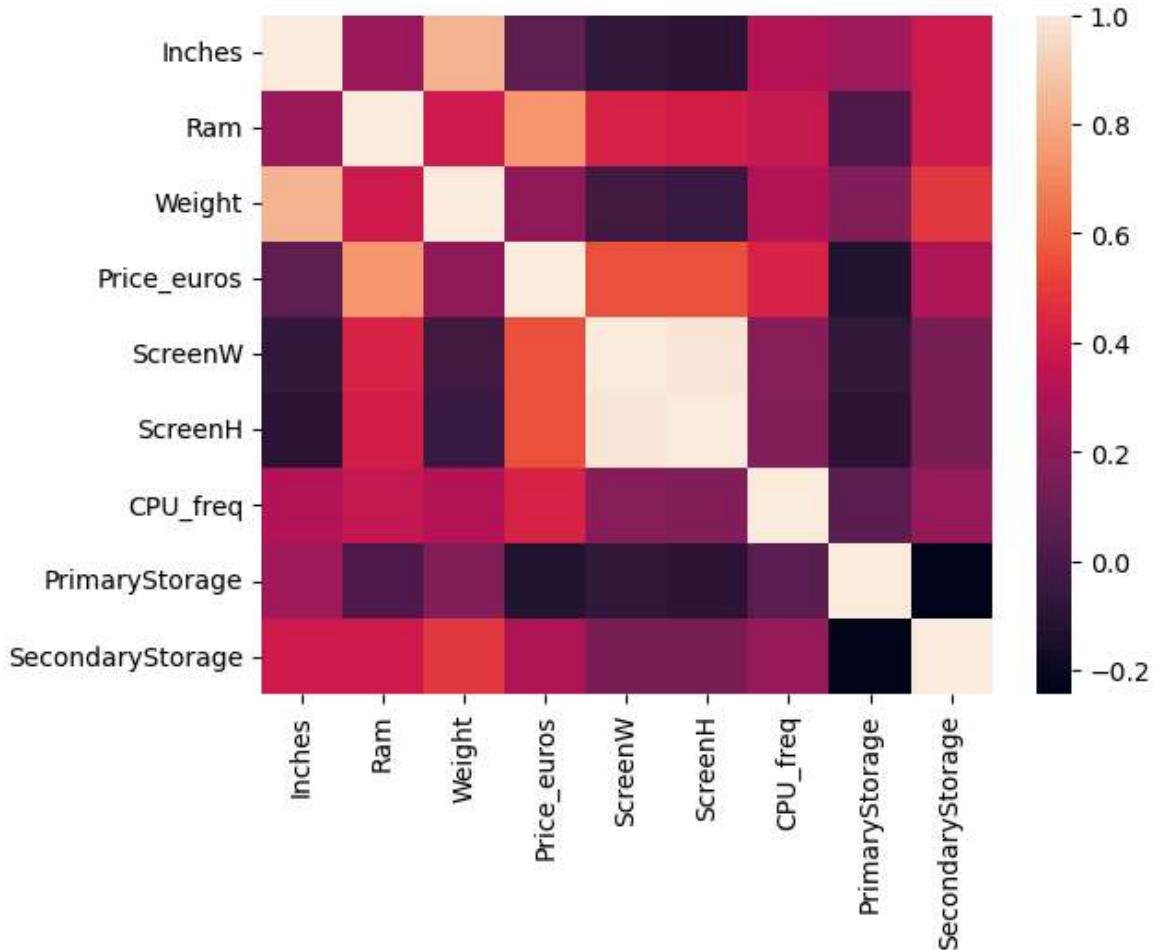
	Inches	Ram	Weight	Price_euros	ScreenW	ScreenH	CPI
<b>Inches</b>	1.000000	0.241078	0.826638	0.066608	-0.068223	-0.093062	0.3
<b>Ram</b>	0.241078	1.000000	0.389370	0.740287	0.424089	0.415241	0.3
<b>Weight</b>	0.826638	0.389370	1.000000	0.211883	-0.028605	-0.050106	0.3
<b>Price_euros</b>	0.066608	0.740287	0.211883	1.000000	0.552491	0.548529	0.4
<b>ScreenW</b>	-0.068223	0.424089	-0.028605	0.552491	1.000000	0.994069	0.1
<b>ScreenH</b>	-0.093062	0.415241	-0.050106	0.548529	0.994069	1.000000	0.1
<b>CPU_freq</b>	0.305037	0.366254	0.318649	0.428847	0.178659	0.164369	1.0
<b>PrimaryStorage</b>	0.264280	0.015365	0.175433	-0.124775	-0.072977	-0.080135	0.0
<b>SecondaryStorage</b>	0.389067	0.390939	0.481495	0.291207	0.146232	0.135293	0.2

```
In [12]: data.corr(numeric_only = True)['Price_euros']
```

```
Out[12]: Inches          0.066608
          Ram            0.740287
          Weight          0.211883
          Price_euros     1.000000
          ScreenW         0.552491
          ScreenH         0.548529
          CPU_freq         0.428847
          PrimaryStorage   -0.124775
          SecondaryStorage 0.291207
          Name: Price_euros, dtype: float64
```

```
In [13]: sns.heatmap(data.corr(numeric_only = True))
```

```
Out[13]: <Axes: >
```



## Exploratory Data Analysis

```
In [15]: data['Company'].value_counts()
```

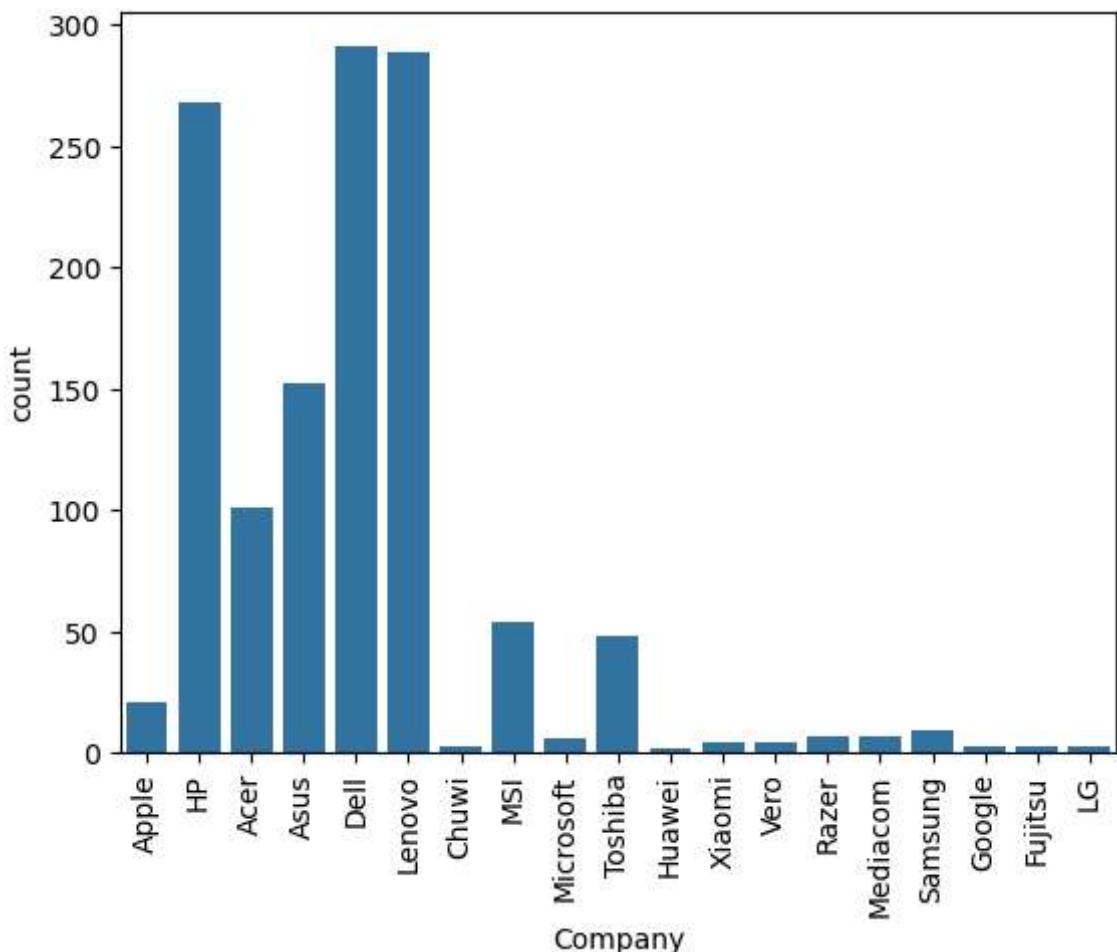
```
Out[15]: Company
Dell      291
Lenovo    289
HP        268
Asus      152
Acer      101
MSI       54
Toshiba   48
Apple     21
Samsung   9
Razer     7
Mediacom  7
Microsoft 6
Xiaomi    4
Vero      4
Chuwi    3
Google    3
Fujitsu   3
LG         3
Huawei    2
Name: count, dtype: int64
```

```
In [16]: # data['Company'].value_counts().plot(kind='bar')

sns.countplot(x='Company', data=data, legend='auto')

# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



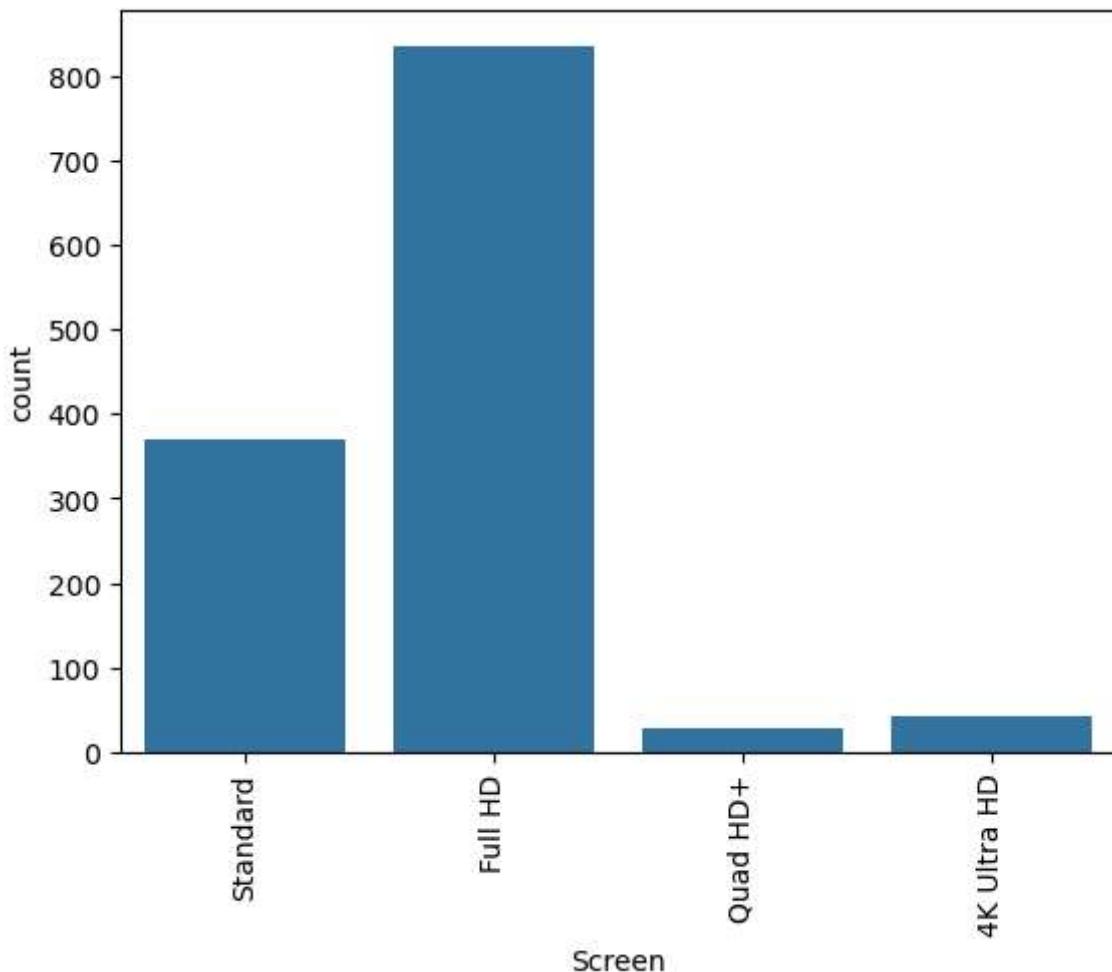
```
In [17]: data['Screen'].value_counts()
```

```
Out[17]: Screen
Full HD      835
Standard     369
4K Ultra HD   43
Quad HD+      28
Name: count, dtype: int64
```

```
In [18]: sns.countplot(x='Screen', data=data, legend='auto')
```

```
# Rotate x-axis Labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



```
In [19]: data['Product'].value_counts()
```

```
Out[19]: Product
XPS 13                30
Inspiron 3567           25
250 G6                 21
Vostro 3568             19
Legion Y520-15IKBN      19
                           ..
VivoBook E201NA          1
Ideapad 520-15IKBR       1
Thinkpad X260             1
Rog G752VL-UH71T          1
X553SA-XX031T (N3050/4GB/500GB/W10) 1
Name: count, Length: 618, dtype: int64
```

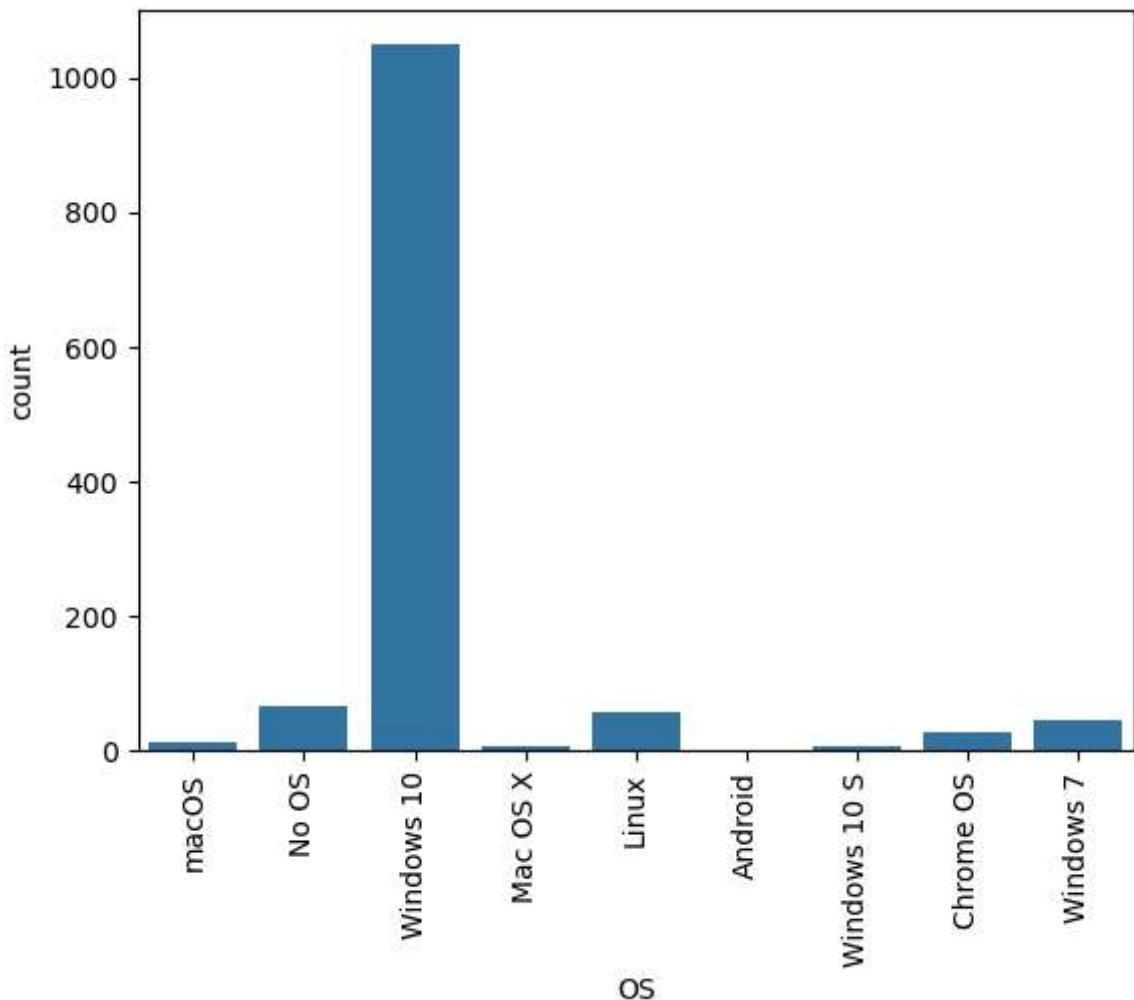
```
In [20]: data['OS'].value_counts()
```

```
Out[20]: OS
Windows 10        1048
No OS              66
Linux               58
Windows 7            45
Chrome OS            27
macOS                13
Mac OS X              8
Windows 10 S            8
Android                2
Name: count, dtype: int64
```

```
In [21]: sns.countplot(x='OS', data=data, legend='auto')

# Rotate x-axis Labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



```
In [22]: data['TypeName'].value_counts()
```

```
Out[22]: TypeName
Notebook          707
Gaming            205
Ultrabook         194
2 in 1 Convertible  117
Workstation        29
Netbook            23
Name: count, dtype: int64
```

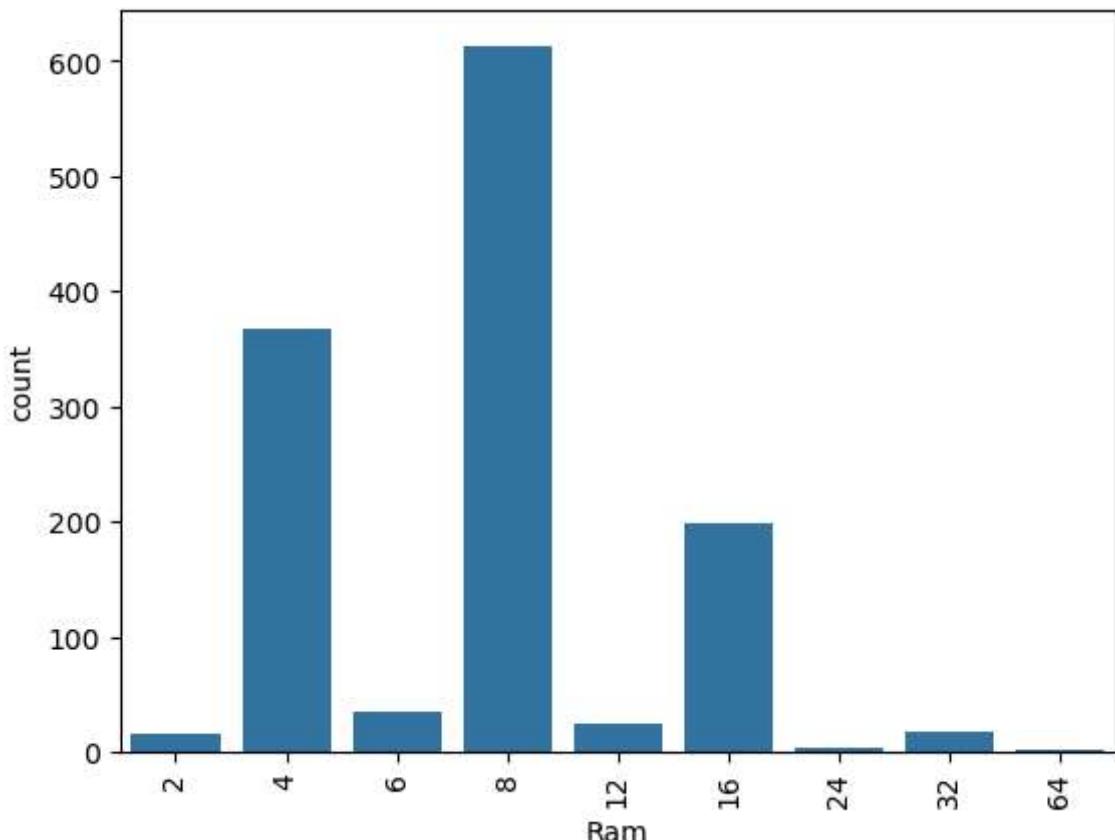
```
In [23]: data['Ram'].value_counts()
```

```
Out[23]: Ram
8      613
4      367
16     198
6      35
12     25
32     17
2      16
24     3
64     1
Name: count, dtype: int64
```

```
In [24]: sns.countplot(x='Ram', data=data, legend='auto')

# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



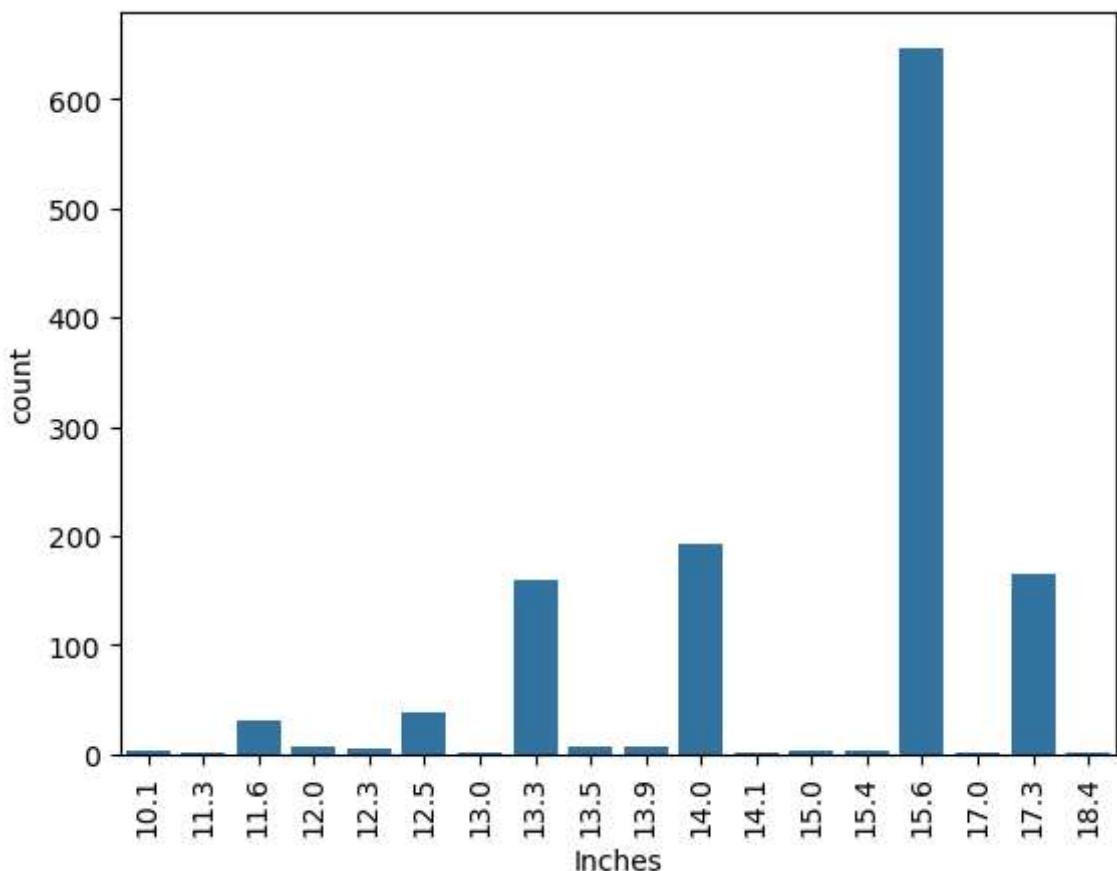
```
In [25]: data['Inches'].value_counts()
```

```
Out[25]: Inches
15.6    647
14.0    193
17.3    164
13.3    160
12.5     39
11.6     31
12.0      6
13.5      6
13.9      6
12.3      5
10.1      4
15.4      4
15.0      4
13.0      2
18.4      1
17.0      1
14.1      1
11.3      1
Name: count, dtype: int64
```

```
In [26]: sns.countplot(x='Inches', data=data, legend='auto')

# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



```
In [27]: data['Touchscreen'].value_counts()
```

```
Out[27]: Touchscreen
No      1087
Yes     188
Name: count, dtype: int64
```

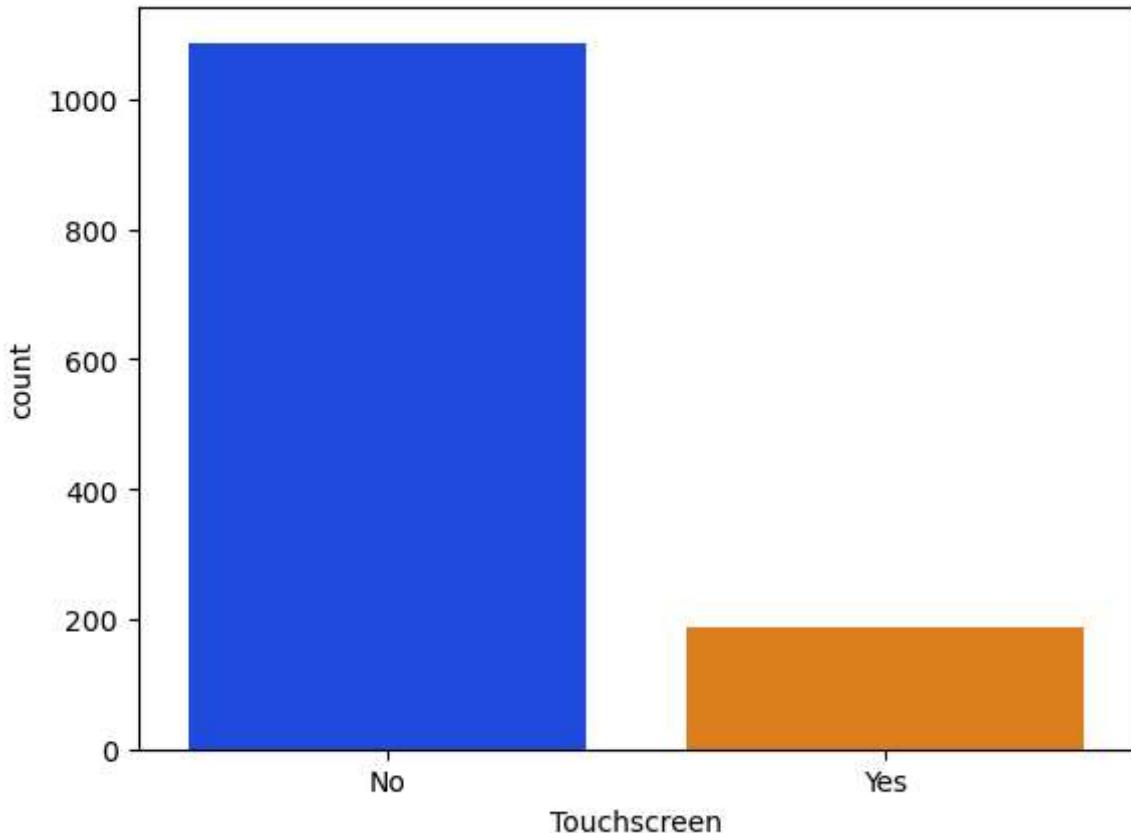
```
In [28]: sns.countplot(x='Touchscreen', data=data, legend='auto', palette='bright' )
```

C:\Users\91700\AppData\Local\Temp\ipykernel\_15640\3302626064.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Touchscreen', data=data, legend='auto', palette='bright' )
```

```
Out[28]: <Axes: xlabel='Touchscreen', ylabel='count'>
```



```
In [29]: data['IPSpanel'].value_counts()
```

```
Out[29]: IPSpanel
No      918
Yes     357
Name: count, dtype: int64
```

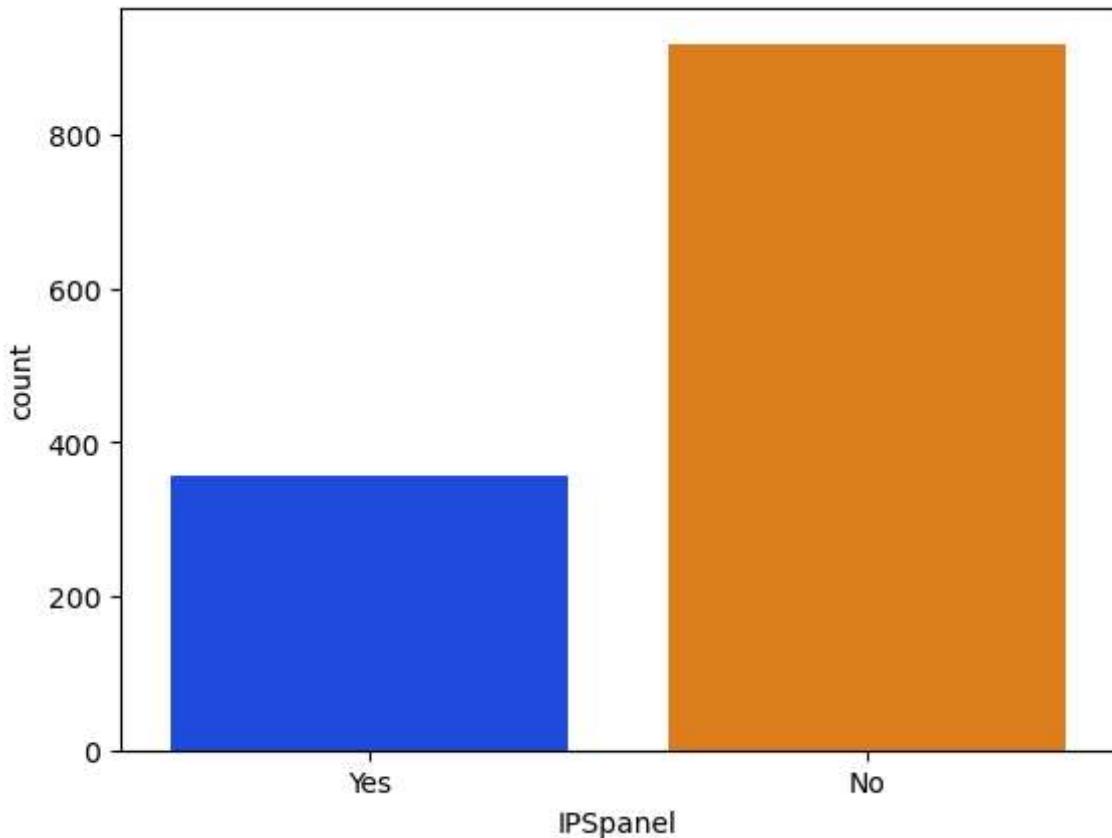
```
In [30]: sns.countplot(x='IPSpanel', data=data, legend='auto', palette='bright' )
```

C:\Users\91700\AppData\Local\Temp\ipykernel\_15640\2524309836.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='IPSpanel', data=data, legend='auto', palette='bright' )
```

```
Out[30]: <Axes: xlabel='IPSpanel', ylabel='count'>
```



```
In [31]: data['RetinaDisplay'].value_counts()
```

```
Out[31]: RetinaDisplay
No      1258
Yes      17
Name: count, dtype: int64
```

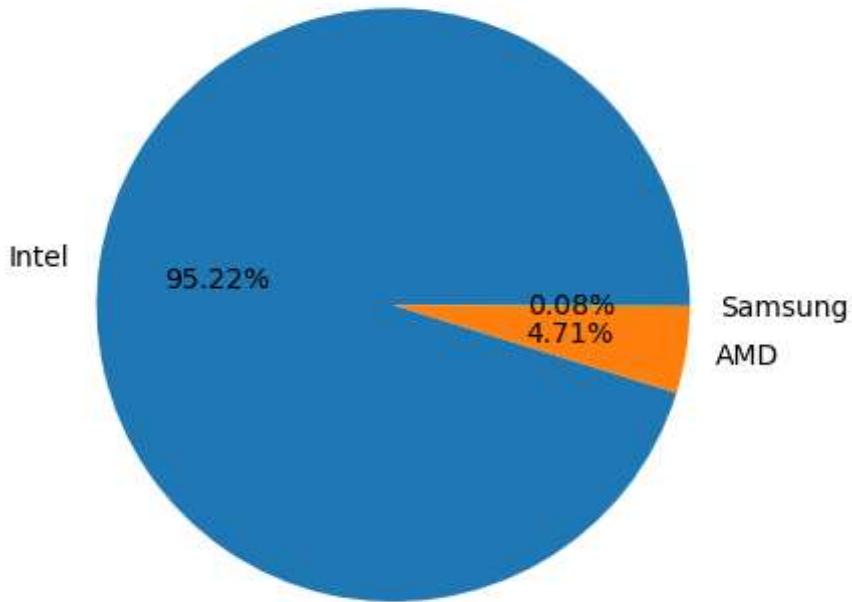
```
In [32]: data['CPU_company'].value_counts()
```

```
Out[32]: CPU_company
Intel      1214
AMD        60
Samsung     1
Name: count, dtype: int64
```

```
In [33]: # Count the occurrences of each CPU_company
CPU_company = data['CPU_company'].value_counts()

# Plot pie chart
plt.pie(CPU_company, labels=CPU_company.index, autopct='%.2f%%')
plt.title("CPU_company")
plt.show()
```

CPU\_company



```
In [34]: data['CPU_model'].value_counts()
```

```
Out[34]: CPU_model
Core i5 7200U      193
Core i7 7700HQ     147
Core i7 7500U      133
Core i3 6006U       81
Core i7 8550U       73
...
Core M m3           1
E-Series E2-9000     1
Core M M3-6Y30       1
A6-Series 7310       1
A9-Series 9410       1
Name: count, Length: 93, dtype: int64
```

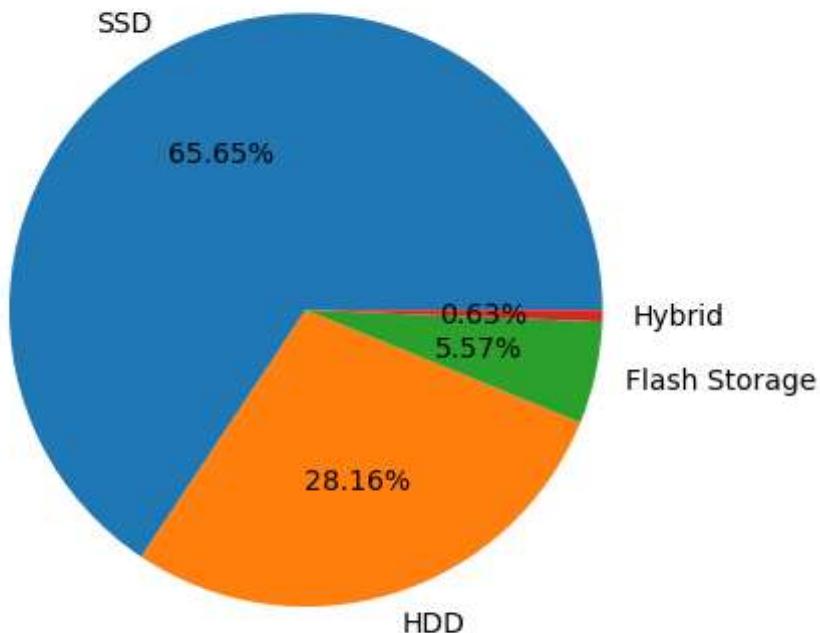
```
In [35]: data['PrimaryStorageType'].value_counts()
```

```
Out[35]: PrimaryStorageType
SSD              837
HDD              359
Flash Storage     71
Hybrid             8
Name: count, dtype: int64
```

```
In [36]: # Count the occurrences of each CPU_company
PrimaryStorageType = data['PrimaryStorageType'].value_counts()

# Plot pie chart
plt.pie(PrimaryStorageType, labels=PrimaryStorageType.index, autopct='%.2f%%')
plt.title("PrimaryStorageType")
plt.show()
```

### PrimaryStorageType



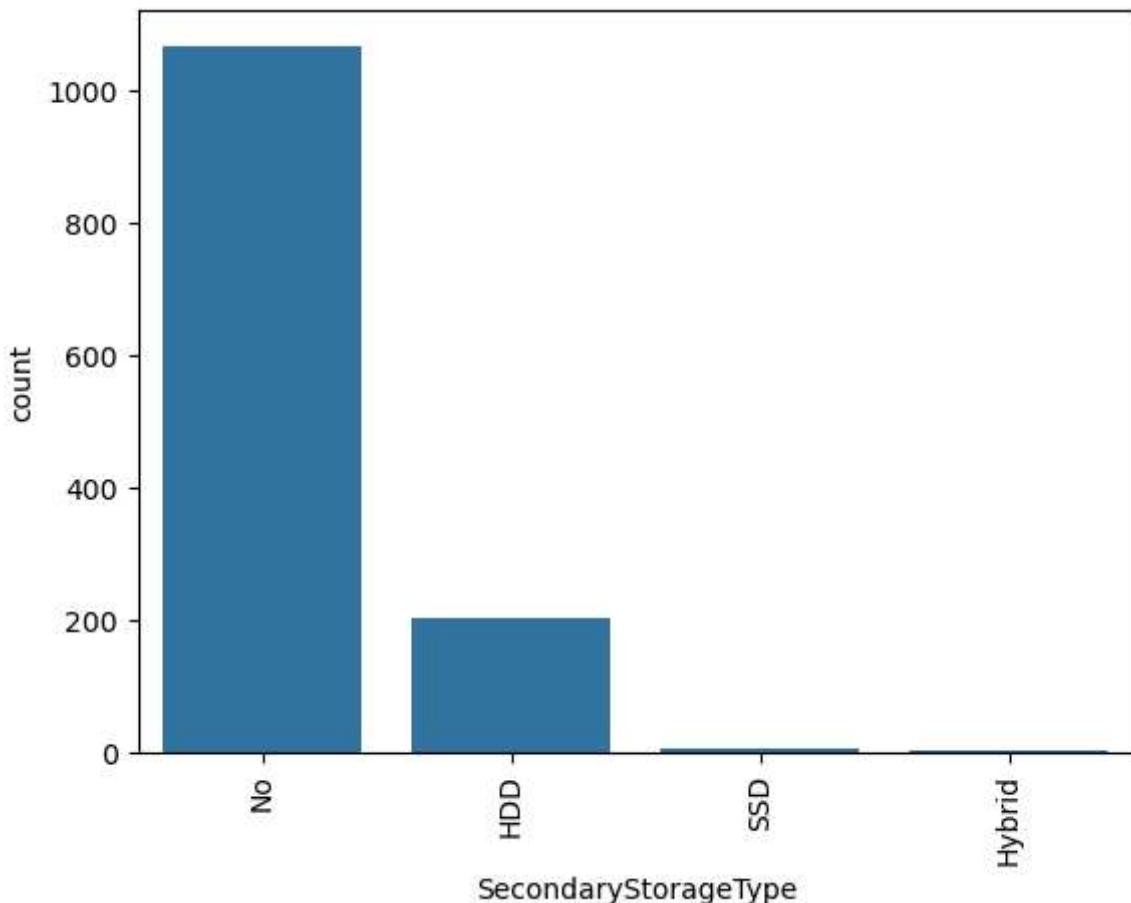
```
In [37]: data['SecondaryStorageType'].value_counts()
```

```
Out[37]: SecondaryStorageType
```

```
No      1067  
HDD     202  
SSD      4  
Hybrid    2  
Name: count, dtype: int64
```

```
In [38]: sns.countplot(x='SecondaryStorageType', data=data, legend='auto')
```

```
# Rotate x-axis labels to be vertical  
plt.xticks(rotation=90)  
  
# Show the plot  
plt.show()
```



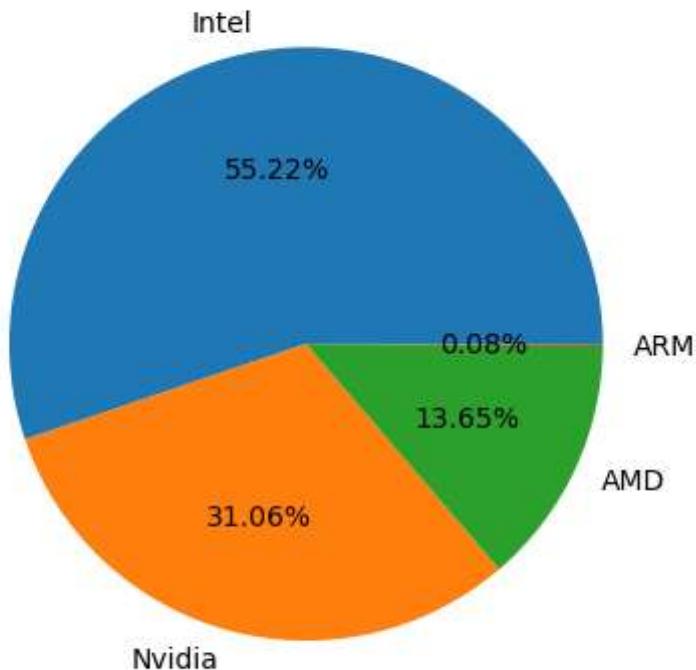
```
In [39]: data['GPU_company'].value_counts()
```

```
Out[39]: GPU_company
Intel      704
Nvidia    396
AMD       174
ARM        1
Name: count, dtype: int64
```

```
In [40]: # Count the occurrences of each CPU_company
GPU_company = data['GPU_company'].value_counts()

# Plot pie chart
plt.pie(GPU_company, labels=GPU_company.index, autopct='%.2f%%')
plt.title("GPU_company")
plt.show()
```

GPU\_company



```
In [41]: data['GPU_model'].value_counts()
```

```
Out[41]: GPU_model
HD Graphics 620      279
HD Graphics 520      181
UHD Graphics 620      68
GeForce GTX 1050      66
GeForce GTX 1060      48
...
Radeon R5 520          1
Radeon R7              1
HD Graphics 540          1
Radeon 540              1
Mali T860 MP4          1
Name: count, Length: 110, dtype: int64
```

```
In [ ]:
```

```
In [43]: # Plot boxplot
sns.boxplot(x="OS", y="Price_euros", data=data, palette="bright")
plt.title("Boxplot of Laptop Prices by Operating System")

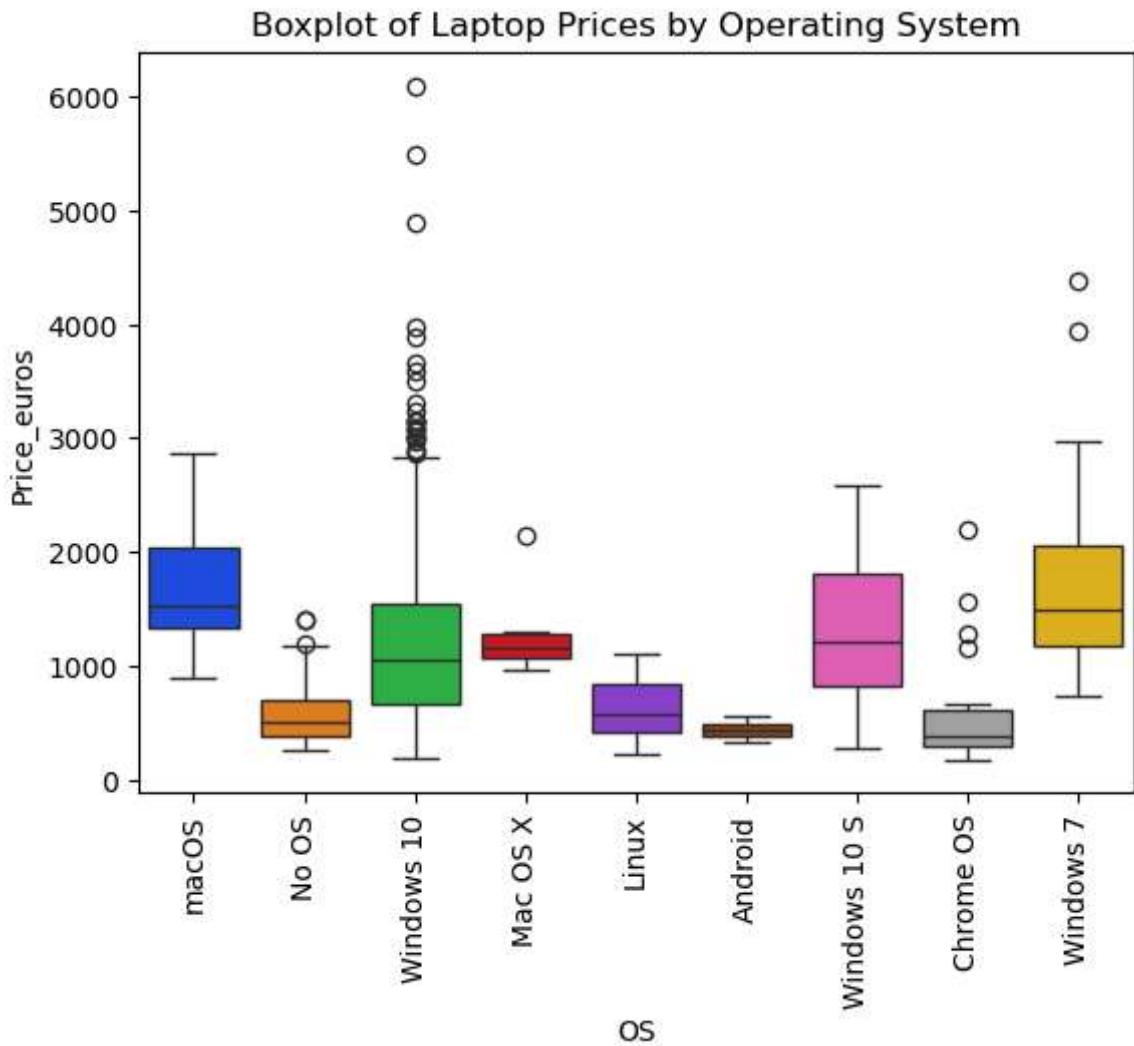
# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```

C:\Users\91700\AppData\Local\Temp\ipykernel\_15640\1734334298.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

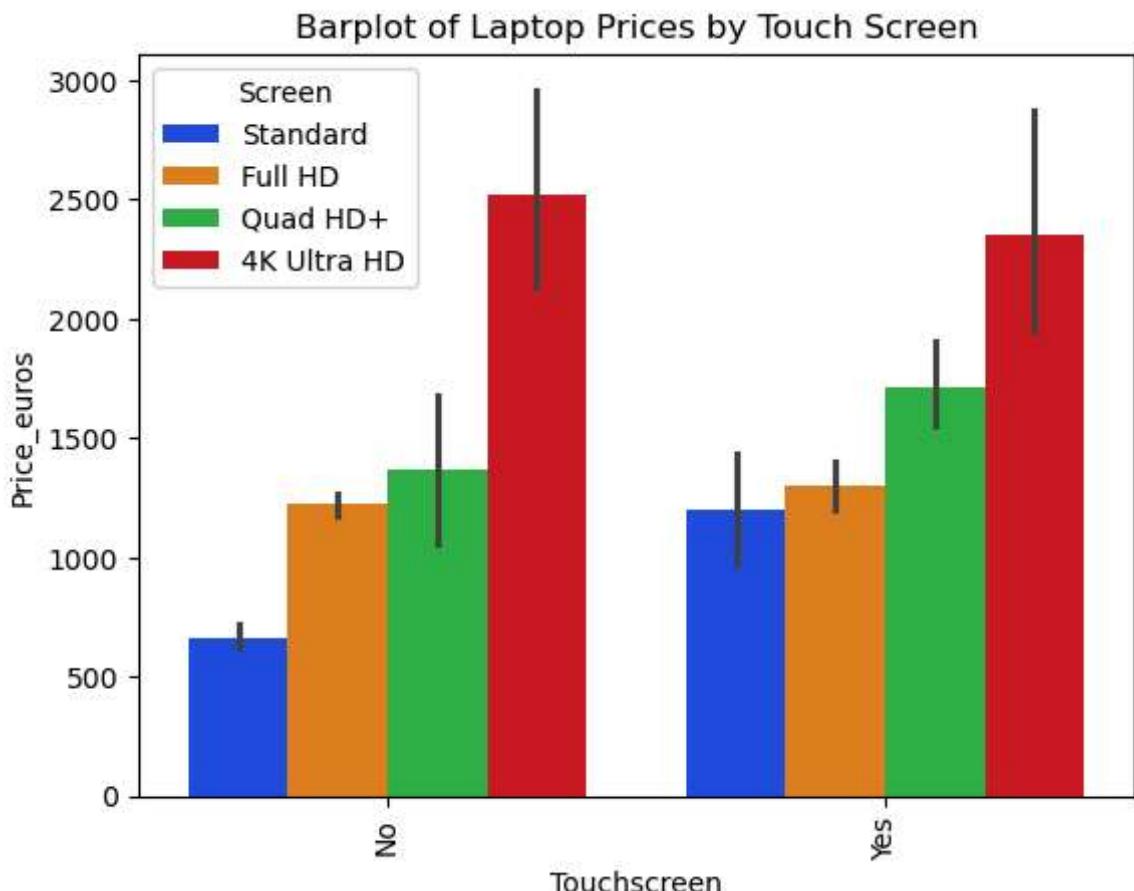
```
sns.boxplot(x="OS", y="Price_euros", data=data, palette="bright")
```



```
In [44]: # Plot barplot
sns.barplot(x = 'Touchscreen', y = 'Price_euros', hue ='Screen', palette="bright"
plt.title("Barplot of Laptop Prices by Touch Screen")

# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```



```
In [45]: # Plot boxplot
sns.barplot(x = 'Touchscreen', y = 'Price_euros', palette="bright", data = data)
plt.title("Barplot of Laptop Prices by Touch Screen")
```

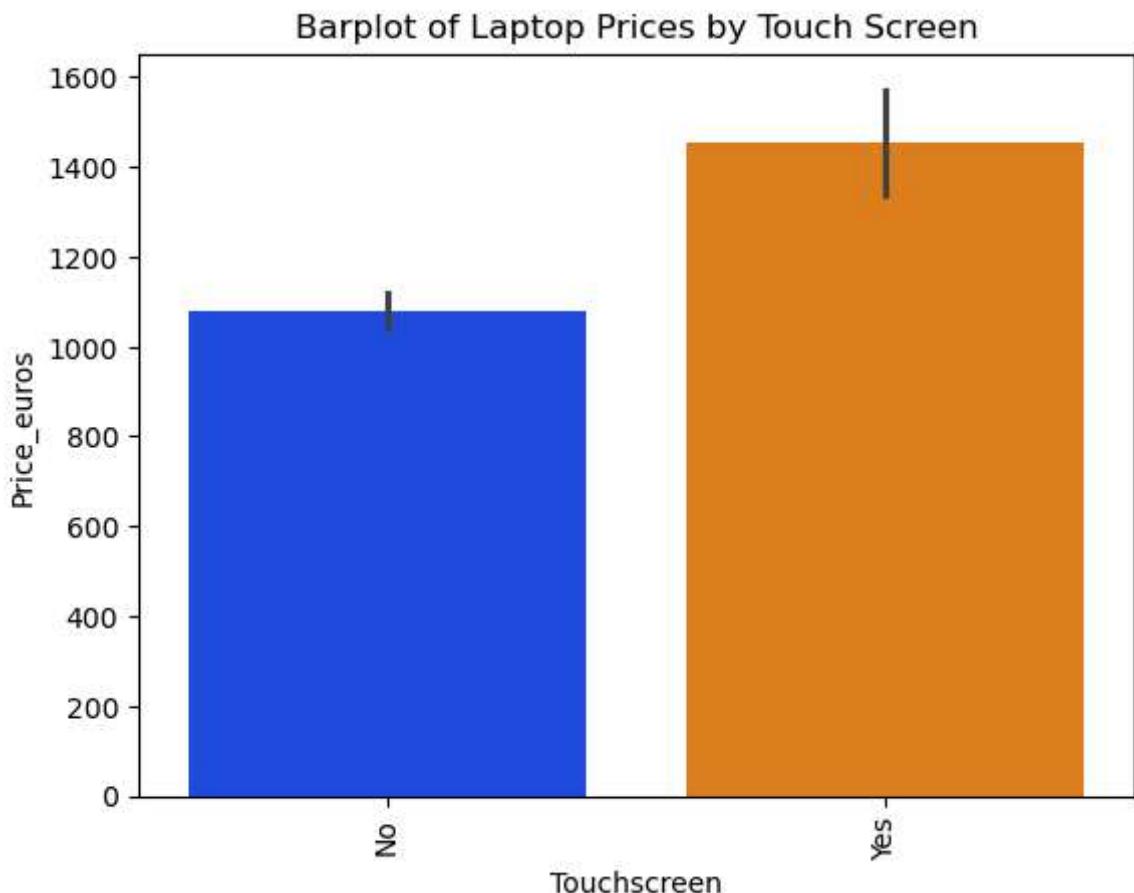
```
# Rotate x-axis Labels to be vertical
plt.xticks(rotation=90)
```

```
# Show the plot
plt.show()
```

C:\Users\91700\AppData\Local\Temp\ipykernel\_15640\3365394172.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x = 'Touchscreen', y = 'Price_euros', palette="bright", data = data)
```

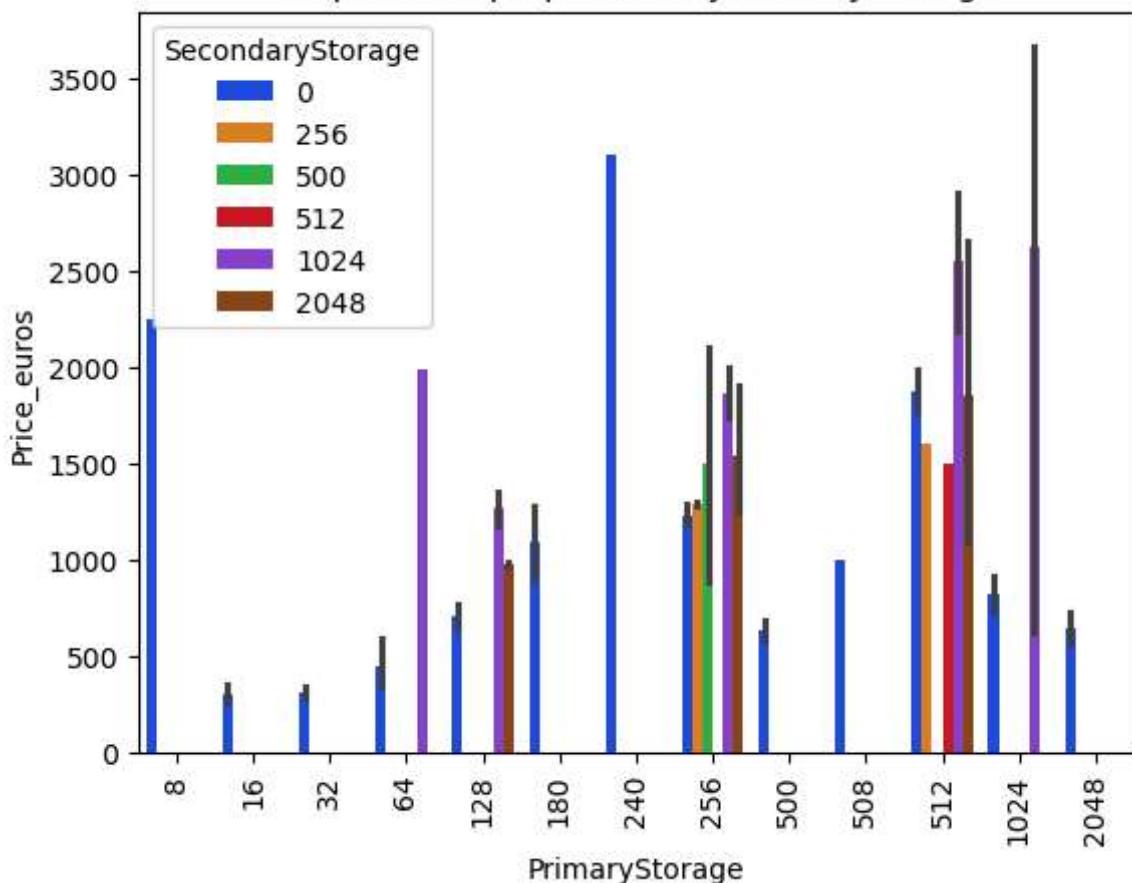


```
In [86]: # Plot boxplot
sns.barplot(x="PrimaryStorage", y="Price_euros", data=data, palette="bright", hue=)
plt.title("Barplot of Laptop Prices by PrimaryStorage")

# Rotate x-axis Labels to be vertical
plt.xticks(rotation=90)

# Show the plot
plt.show()
```

## Barplot of Laptop Prices by PrimaryStorage

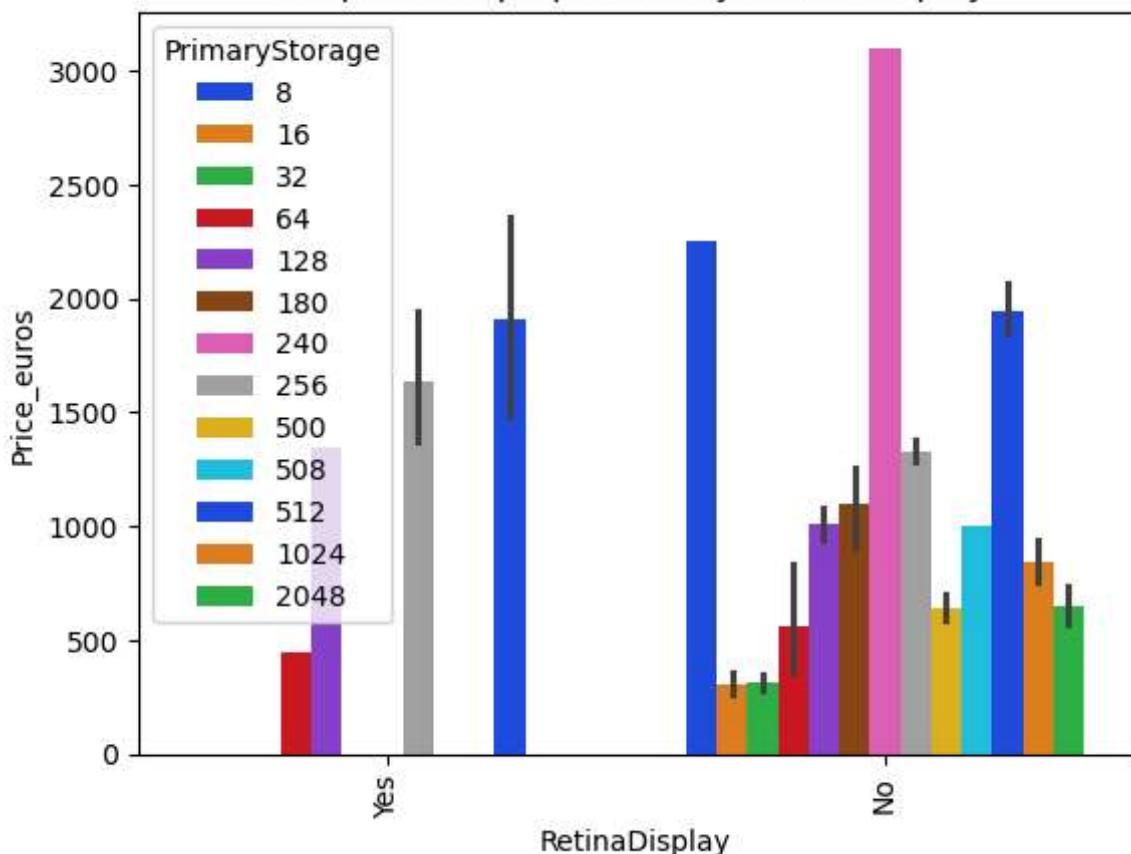


```
In [88]: # Plot boxplot
sns.barplot(x="RetinaDisplay", y="Price_euros", data=data, palette="bright", hue=RetinaDisplay)
plt.title("Barplot of Laptop Prices by Retina Display")

# Rotate x-axis labels to be vertical
plt.xticks(rotation=90)

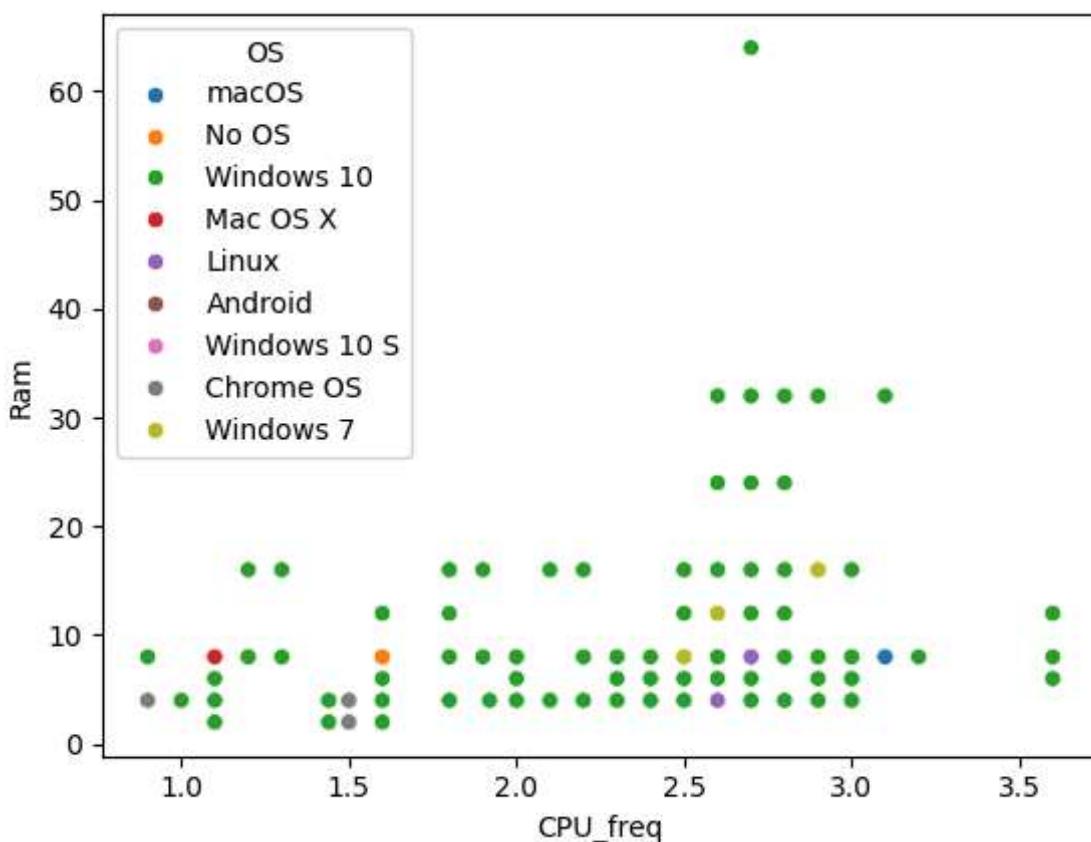
# Show the plot
plt.show()
```

### Barplot of Laptop Prices by Retina Display



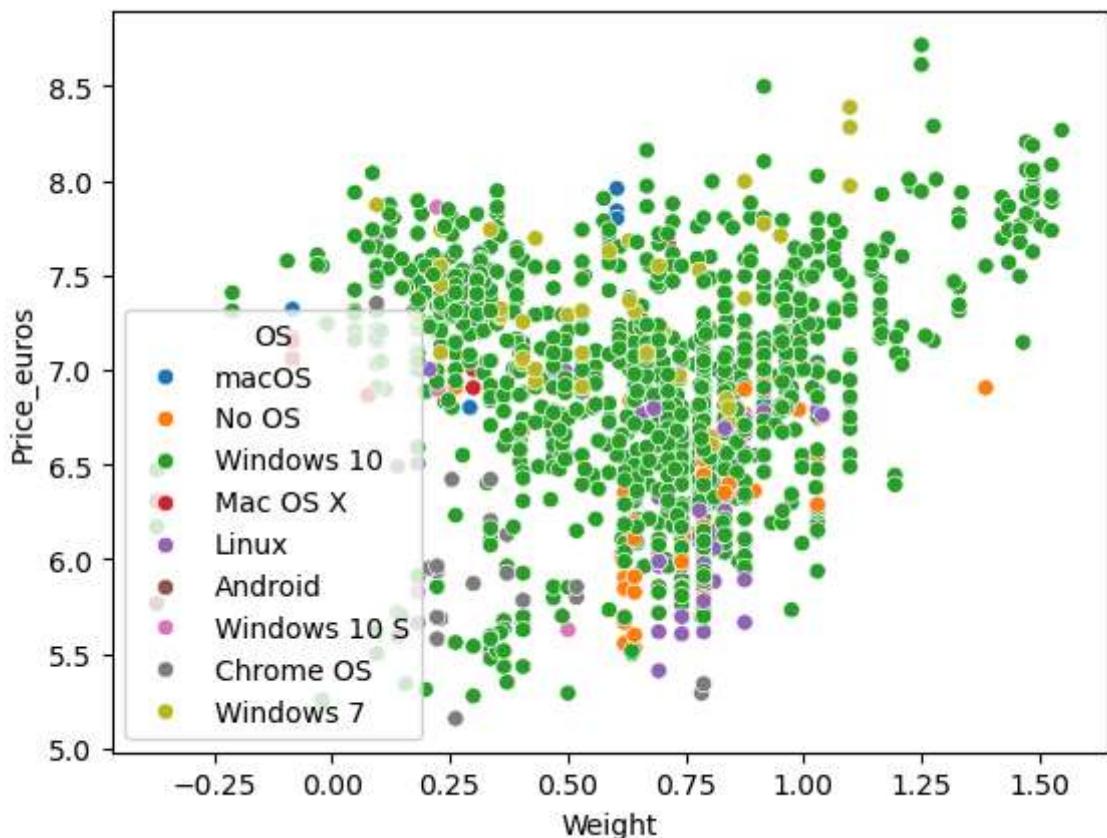
```
In [90]: sns.scatterplot(data = data , x= 'CPU_freq', y = 'Ram', hue = 'OS')
```

```
Out[90]: <Axes: xlabel='CPU_freq', ylabel='Ram'>
```



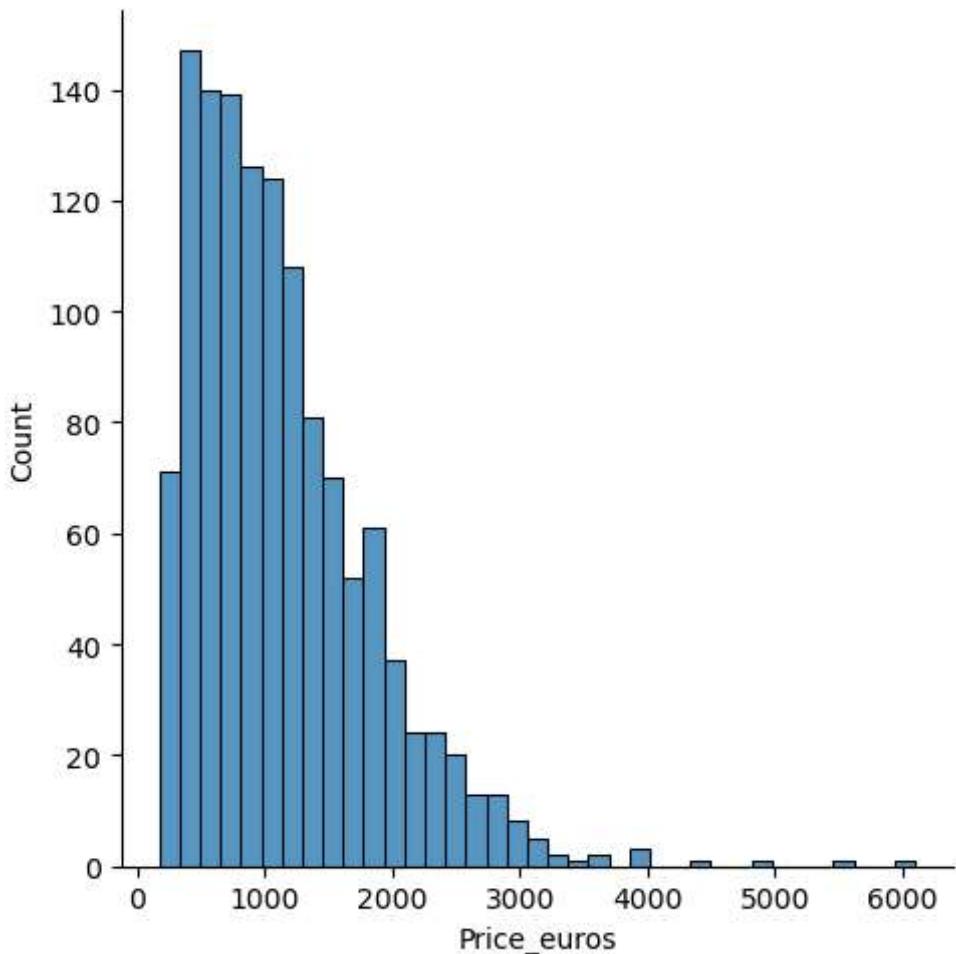
```
In [92]: sns.scatterplot(data = data , x= np.log(data['Weight']), y = np.log(data['Price']))
```

```
Out[92]: <Axes: xlabel='Weight', ylabel='Price_euros'>
```



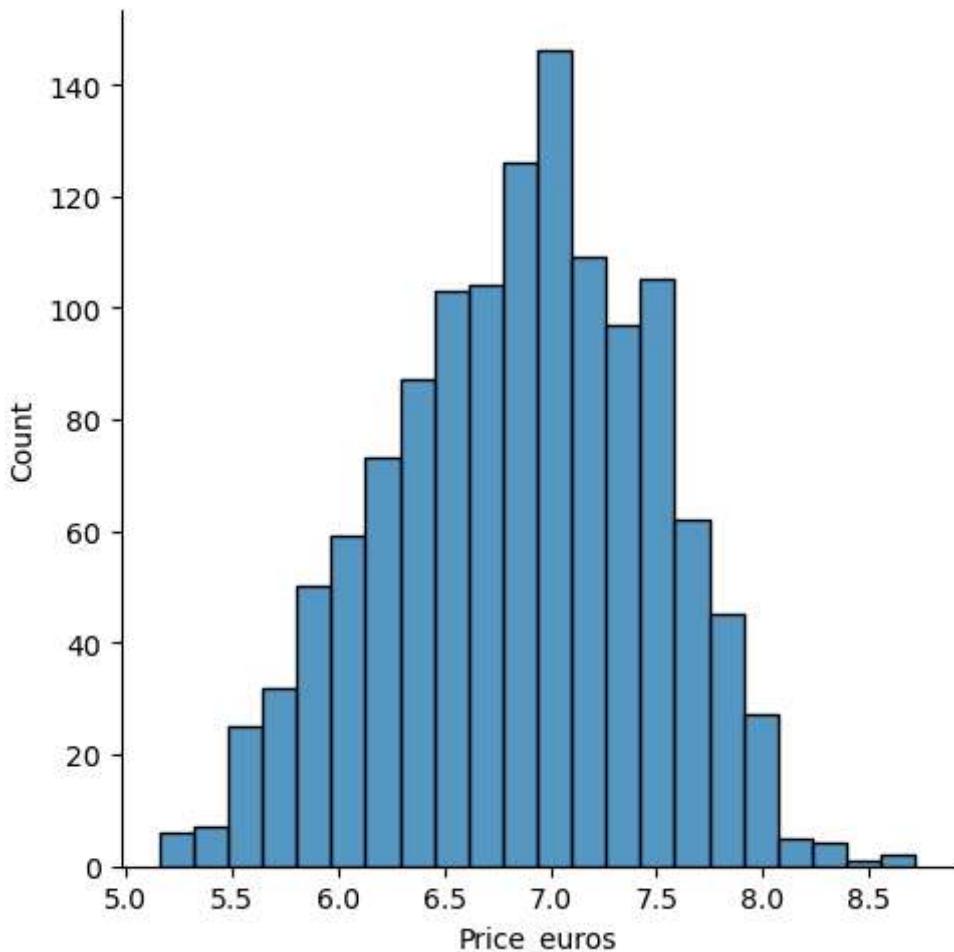
```
In [94]: sns.displot(data['Price_euros'])
```

```
Out[94]: <seaborn.axisgrid.FacetGrid at 0x28a02a40950>
```



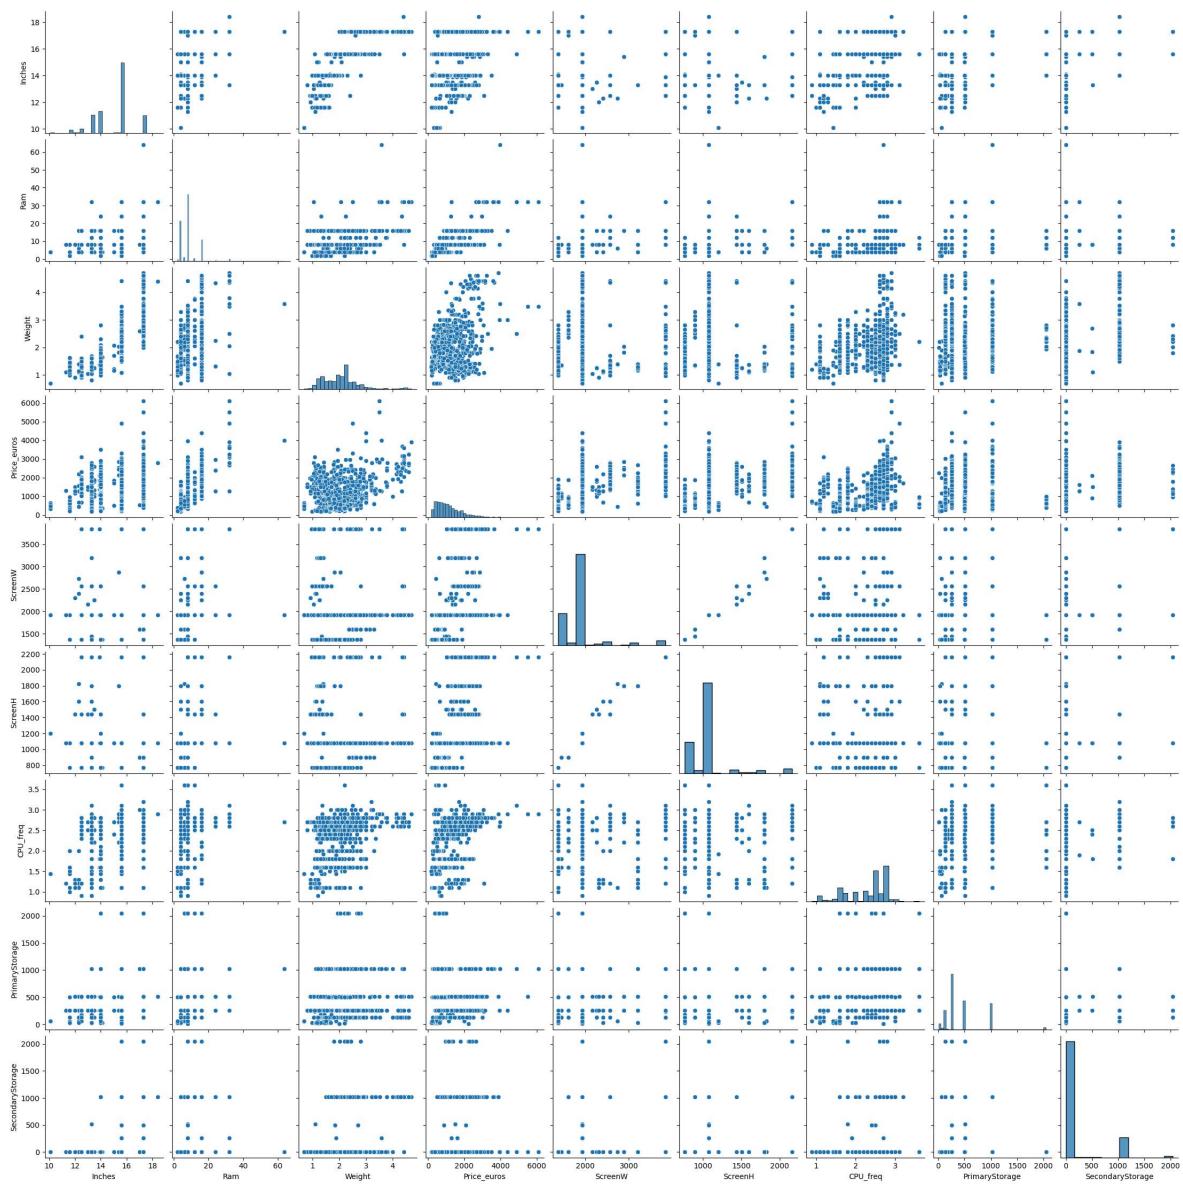
```
In [96]: sns.displot(x= np.log(data['Price_euros']))
```

```
Out[96]: <seaborn.axisgrid.FacetGrid at 0x28a094f6390>
```



```
In [100]: sns.pairplot(data)
```

```
Out[100]: <seaborn.axisgrid.PairGrid at 0x28a03838da0>
```



## Define features and target

```
In [102...]: X = data.drop('Price_euros', axis=1)
y = np.log(data['Price_euros'])
```

## Identify categorical and numerical columns

```
In [104...]: categorical_cols = X.select_dtypes(include=['object']).columns
categorical_cols
```

```
Out[104...]: Index(['Company', 'Product', 'TypeName', 'OS', 'Screen', 'Touchscreen',
       'IPSpanel', 'RetinaDisplay', 'CPU_company', 'CPU_model',
       'PrimaryStorageType', 'SecondaryStorageType', 'GPU_company',
       'GPU_model'],
      dtype='object')
```

```
In [105...]: numerical_cols = X.select_dtypes(include=['float64', 'int64', 'int32']).columns
numerical_cols
```

```
Out[105...]: Index(['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'CPU_freq',
       'PrimaryStorage', 'SecondaryStorage'],
      dtype='object')
```

## Split the data

```
In [107...]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

## Encoding for categorical features and Preprocessing pipelines

```
In [109...]: preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), [
            'Company', 'Touchscreen', 'CPU_core', 'SecondaryStorage']),
        ('num', StandardScaler(), [
            'Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'CPU_freq']),
        ('rest', 'passthrough', [
            'SecondaryStorage'])],
    remainder = 'passthrough')

regressor = LinearRegression()

pipe_reg = Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', regressor)
])

# Train the Model
pipe_reg.fit(X_train, y_train)

y_pred_reg = pipe_reg.predict(X_test)

# Evaluate the test model
mse_test_reg = mean_squared_error(y_test, y_pred_reg)
r2_test_reg = r2_score(y_test, y_pred_reg)

print(f"Mean Squared Error Test data: {mse_test_reg}")
print(f"R-squared Test data: {r2_test_reg}")

Mean Squared Error Test data: 0.04296617417476588
R-squared Test data: 0.8749081962913069
D:\Anaconda\Lib\site-packages\sklearn\preprocessing\_encoders.py:242: UserWarning:
g: Found unknown categories in columns [1, 9, 13] during transform. These unknown
categories will be encoded as all zeros
warnings.warn(
```

## Ridge Regression

```
In [112...]: rid_prepro = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), [
            'Company', 'Touchscreen', 'CPU_core', 'SecondaryStorage']),
        ('rest', 'passthrough', [
            'SecondaryStorage'])],
    remainder = 'passthrough')
```

```

('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH',
                           'SecondaryStorage'])
], remainder = 'passthrough')

rid_reg = Ridge(alpha=0.9)

pipe_rid = Pipeline([
    ('rid_prep', rid_prep),
    ('rid_reg', rid_reg)
])

# Train the Model
pipe_rid.fit(X_train, y_train)

y_pred_rid = pipe_rid.predict(X_test)

# Evaluate the test model
mse_test_rid = mean_squared_error(y_test, y_pred_rid)
r2_test_rid = r2_score(y_test, y_pred_rid)

print(f"Mean Squared Error Test data: {mse_test_rid}")
print(f"R-squared Test data: {r2_test_rid}")

```

Mean Squared Error Test data: 0.030172629390433903  
R-squared Test data: 0.9121553476525265

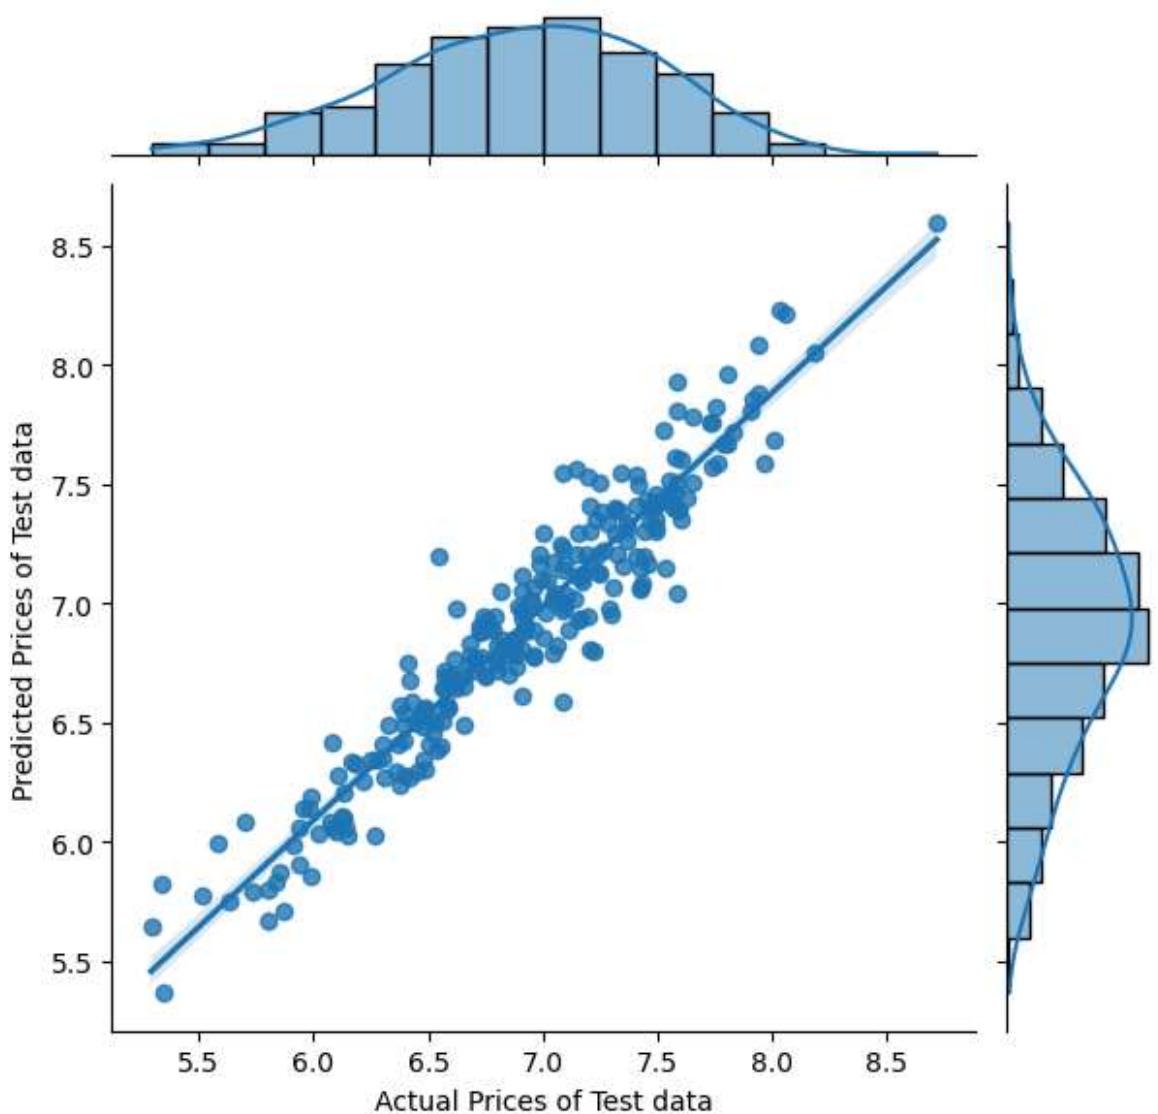
D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning:  
g: Found unknown categories in columns [1, 9, 13] during transform. These unknown  
categories will be encoded as all zeros  
warnings.warn(

In [165...]

```

# Visualize test results
sns.jointplot(x = y_test, y= y_pred_rid, kind="reg")
plt.xlabel("Actual Prices of Test data")
plt.ylabel("Predicted Prices of Test data")
plt.show()

```



## Lasso Regression

```
In [115]: las_prep = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'TouchScreen', 'CPU_core', 'SecondaryStorage'])
    ], remainder = 'passthrough')

las_reg = Lasso(alpha=0.0001)

pipe_las = Pipeline([
    ('las_prep', las_prep),
    ('las_reg', las_reg)
])

# Train the Model
pipe_las.fit(X_train, y_train)

y_pred_las = pipe_las.predict(X_test)
```

```
# Evaluate the test model
mse_test_las = mean_squared_error(y_test, y_pred_las)
r2_test_las = r2_score(y_test, y_pred_las)

print(f"Mean Squared Error Test data: {mse_test_las}")
print(f"R-squared Test data: {r2_test_las}")
```

Mean Squared Error Test data: 0.0312409212885543  
R-squared Test data: 0.9090451205264226

D:\Anaconda\Lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:658: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 2.0202187520186685, tolerance: 0.040076899266623674  
model = cd\_fast.sparse\_enet\_coordinate\_descent(  
D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning: Found unknown categories in columns [1, 9, 13] during transform. These unknown categories will be encoded as all zeros  
warnings.warn(

## KNeighbors Regressor

In [118...]

```
knn_prepro = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenSecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'Touchscreen', 'CPU_core', 'ScreenResolution']),
    ], remainder = 'passthrough')

knn = KNeighborsRegressor(n_neighbors=5)

pipe_knn = Pipeline([
    ('knn_prepro', knn_prepro),
    ('knn', knn)
])

# Train the Model
pipe_knn.fit(X_train, y_train)

y_pred_knn = pipe_knn.predict(X_test)

# Evaluate the test model
mse_test_knn = mean_squared_error(y_test, y_pred_knn)
r2_test_knn = r2_score(y_test, y_pred_knn)

print(f"Mean Squared Error Test data: {mse_test_knn}")
print(f"R-squared Test data: {r2_test_knn}")
```

Mean Squared Error Test data: 0.05207344800605449  
R-squared Test data: 0.84839326140809

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning: Found unknown categories in columns [1, 9, 13] during transform. These unknown categories will be encoded as all zeros  
warnings.warn(

## Decision Tree Regressor

```
In [121...]: dtr_prep = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'Touchscreen', 'CPU_core', 'ScreenType']),
    ], remainder = 'passthrough')

dtr = DecisionTreeRegressor(max_depth=12, min_samples_split=12, max_features=0.8)

pipe_dtr = Pipeline([
    ('dtr_prep', dtr_prep),
    ('dtr', dtr)
])

# Train the Model
pipe_dtr.fit(X_train, y_train)

y_pred_dtr = pipe_dtr.predict(X_test)

# Evaluate the test model
mse_test_dtr = mean_squared_error(y_test, y_pred_dtr)
r2_test_dtr = r2_score(y_test, y_pred_dtr)

print(f"Mean Squared Error Test data: {mse_test_dtr}")
print(f"R-squared Test data: {r2_test_dtr}")
```

Mean Squared Error Test data: 0.06291982229794918  
 R-squared Test data: 0.8168151060351233

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning:  
 g: Found unknown categories in columns [1, 9, 13] during transform. These unknown  
 categories will be encoded as all zeros  
 warnings.warn(

## Random Forest Regressor

```
In [124...]: rfr_prep = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'Touchscreen', 'CPU_core', 'ScreenType']),
    ], remainder = 'passthrough')

rfr = RandomForestRegressor(n_estimators = 100, max_features=0.5, max_depth=15, n_jobs=-1)

pipe_rfr = Pipeline([
    ('rfr_prep', rfr_prep),
    ('rfr', rfr)
])

# Train the Model
pipe_rfr.fit(X_train, y_train)
```

```

y_pred_rfr = pipe_rfr.predict(X_test)

# Evaluate the test model
mse_test_rfr = mean_squared_error(y_test, y_pred_rfr)
r2_test_rfr = r2_score(y_test, y_pred_rfr)

print(f"Mean Squared Error Test data: {mse_test_rfr}")
print(f"R-squared Test data: {r2_test_rfr}")

```

Mean Squared Error Test data: 0.03839887502774533  
 R-squared Test data: 0.8882054399801229

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning:  
 g: Found unknown categories in columns [1, 9, 13] during transform. These unknown  
 categories will be encoded as all zeros  
 warnings.warn(

## AdaBoost Regressor

```

In [127]: ab_prepro = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH',
                                    'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company',
                            'TouchScreen', 'CPU_ClockSpeed', 'RAM'],
         remainder = 'passthrough')

ab_regg = AdaBoostRegressor(n_estimators =500, learning_rate=0.9, random_state=3)

pipe_ab = Pipeline([
    ('ab_prepro', ab_prepro),
    ('ab_regg', ab_regg)
])

# Train the Model
pipe_ab.fit(X_train, y_train)

y_pred_ab = pipe_ab.predict(X_test)

# Evaluate the test model
mse_test_ab = mean_squared_error(y_test, y_pred_ab)
r2_test_ab = r2_score(y_test, y_pred_ab)

print(f"Mean Squared Error Test data: {mse_test_ab}")
print(f"R-squared Test data: {r2_test_ab}")

```

Mean Squared Error Test data: 0.06575269858865244  
 R-squared Test data: 0.8085674644497624

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning:  
 g: Found unknown categories in columns [1, 9, 13] during transform. These unknown  
 categories will be encoded as all zeros  
 warnings.warn(

## Gradient Boosting Regressor

```
In [130...]: gbr_repro = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'Touchscreen', 'CPU_core', 'SecondaryStorage']),
    ], remainder = 'passthrough')

gb_regg = GradientBoostingRegressor(n_estimators = 1000, random_state=3, alpha=0.001)

pipe_gb = Pipeline([
    ('gbr_repro', gbr_repro),
    ('gb_regg', gb_regg)
])

# Train the Model
pipe_gb.fit(X_train, y_train)

y_pred_gb = pipe_gb.predict(X_test)

# Evaluate the test model
mse_test_gb = mean_squared_error(y_test, y_pred_gb)
r2_test_gb = r2_score(y_test, y_pred_gb)

print(f"Mean Squared Error Test data: {mse_test_gb}")
print(f"R-squared Test data: {r2_test_gb}")
```

Mean Squared Error Test data: 0.03388778820575893  
 R-squared Test data: 0.901339026995655

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning:  
 g: Found unknown categories in columns [1, 9, 13] during transform. These unknown  
 categories will be encoded as all zeros  
 warnings.warn(

## XGBoost

```
In [133...]: xgb_repro = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Inches', 'Ram', 'Weight', 'ScreenW', 'ScreenH', 'SecondaryStorage']),
        ('cat', OneHotEncoder(drop = 'first', handle_unknown= 'ignore'), ['Company', 'Touchscreen', 'CPU_core', 'SecondaryStorage']),
    ], remainder = 'passthrough')

xgb_regg = XGBRegressor(n_estimators = 100, learning_rate=0.9, max_depth=10, random_state=3)

pipe_xgb = Pipeline([
    ('xgb_repro', xgb_repro),
    ('xgb_regg', xgb_regg)
])

# Train the Model
pipe_xgb.fit(X_train, y_train)
```

```
y_pred_xgb = pipe_xgb.predict(X_test)

# Evaluate the test model
mse_test_xgb = mean_squared_error(y_test, y_pred_xgb)
r2_test_xgb = r2_score(y_test, y_pred_xgb)

print(f"Mean Squared Error Test data: {mse_test_xgb}")
print(f"R-squared Test data: {r2_test_xgb}")
```

Mean Squared Error Test data: 0.05376753239030201

R-squared Test data: 0.8434611008112847

D:\Anaconda\Lib\site-packages\sklearn\preprocessing\\_encoders.py:242: UserWarning: Found unknown categories in columns [1, 9, 13] during transform. These unknown categories will be encoded as all zeros

```
warnings.warn(
```

## Exporting the Model

In [137...]

```
import pickle

# pickle.dump(pipe_rid,open('model.pkl','wb'))

# Streamlit-notworking
pickle.dump(data,open('data.pkl','wb'))
pickle.dump(pipe_rid,open('pipe.pkl','wb'))
```